

## 1.Importing Necessary Libraries

```
In [1]: import pandas as pd  
from matplotlib import pyplot as plt  
import seaborn as sns
```

## 2. Importing data



```
In [2]: data_startups = pd.read_csv('50_Startups.csv')
data_startups
```

```
Out[2]:
```

|    | R&D Spend | Administration | Marketing Spend | State      | Profit    |
|----|-----------|----------------|-----------------|------------|-----------|
| 0  | 165349.20 | 136897.80      | 471784.10       | New York   | 192261.83 |
| 1  | 162597.70 | 151377.59      | 443898.53       | California | 191792.06 |
| 2  | 153441.51 | 101145.55      | 407934.54       | Florida    | 191050.39 |
| 3  | 144372.41 | 118671.85      | 383199.62       | New York   | 182901.99 |
| 4  | 142107.34 | 91391.77       | 366168.42       | Florida    | 166187.94 |
| 5  | 131876.90 | 99814.71       | 362861.36       | New York   | 156991.12 |
| 6  | 134615.46 | 147198.87      | 127716.82       | California | 156122.51 |
| 7  | 130298.13 | 145530.06      | 323876.68       | Florida    | 155752.60 |
| 8  | 120542.52 | 148718.95      | 311613.29       | New York   | 152211.77 |
| 9  | 123334.88 | 108679.17      | 304981.62       | California | 149759.96 |
| 10 | 101913.08 | 110594.11      | 229160.95       | Florida    | 146121.95 |
| 11 | 100671.96 | 91790.61       | 249744.55       | California | 144259.40 |
| 12 | 93863.75  | 127320.38      | 249839.44       | Florida    | 141585.52 |
| 13 | 91992.39  | 135495.07      | 252664.93       | California | 134307.35 |
| 14 | 119943.24 | 156547.42      | 256512.92       | Florida    | 132602.65 |
| 15 | 114523.61 | 122616.84      | 261776.23       | New York   | 129917.04 |
| 16 | 78013.11  | 121597.55      | 264346.06       | California | 126992.93 |
| 17 | 94657.16  | 145077.58      | 282574.31       | New York   | 125370.37 |
| 18 | 91749.16  | 114175.79      | 294919.57       | Florida    | 124266.90 |
| 19 | 86419.70  | 153514.11      | 0.00            | New York   | 122776.86 |
| 20 | 76253.86  | 113867.30      | 298664.47       | California | 118474.03 |
| 21 | 78389.47  | 153773.43      | 299737.29       | New York   | 111313.02 |
| 22 | 73994.56  | 122782.75      | 303319.26       | Florida    | 110352.25 |
| 23 | 67532.53  | 105751.03      | 304768.73       | Florida    | 108733.99 |
| 24 | 77044.01  | 99281.34       | 140574.81       | New York   | 108552.04 |
| 25 | 64664.71  | 139553.16      | 137962.62       | California | 107404.34 |
| 26 | 75328.87  | 144135.98      | 134050.07       | Florida    | 105733.54 |
| 27 | 72107.60  | 127864.55      | 353183.81       | New York   | 105008.31 |
| 28 | 66051.52  | 182645.56      | 118148.20       | Florida    | 103282.38 |
| 29 | 65605.48  | 153032.06      | 107138.38       | New York   | 101004.64 |
| 30 | 61994.48  | 115641.28      | 91131.24        | Florida    | 99937.59  |
| 31 | 61136.38  | 152701.92      | 88218.23        | New York   | 97483.56  |
| 32 | 63408.86  | 129219.61      | 46085.25        | California | 97427.84  |
| 33 | 55493.95  | 103057.49      | 214634.81       | Florida    | 96778.92  |

|    | R&D Spend | Administration | Marketing Spend | State      | Profit   |
|----|-----------|----------------|-----------------|------------|----------|
| 34 | 46426.07  | 157693.92      | 210797.67       | California | 96712.80 |
| 35 | 46014.02  | 85047.44       | 205517.64       | New York   | 96479.51 |
| 36 | 28663.76  | 127056.21      | 201126.82       | Florida    | 90708.19 |
| 37 | 44069.95  | 51283.14       | 197029.42       | California | 89949.14 |
| 38 | 20229.59  | 65947.93       | 185265.10       | New York   | 81229.06 |
| 39 | 38558.51  | 82982.09       | 174999.30       | California | 81005.76 |
| 40 | 28754.33  | 118546.05      | 172795.67       | California | 78239.91 |
| 41 | 27892.92  | 84710.77       | 164470.71       | Florida    | 77798.83 |
| 42 | 23640.93  | 96189.63       | 148001.11       | California | 71498.49 |
| 43 | 15505.73  | 127382.30      | 35534.17        | New York   | 69758.98 |
| 44 | 22177.74  | 154806.14      | 28334.72        | California | 65200.33 |
| 45 | 1000.23   | 124153.04      | 1903.93         | New York   | 64926.08 |
| 46 | 1315.46   | 115816.21      | 297114.46       | Florida    | 49490.75 |
| 47 | 0.00      | 135426.92      | 0.00            | California | 42559.73 |
| 48 | 542.05    | 51743.15       | 0.00            | New York   | 35673.41 |
| 49 | 0.00      | 116983.80      | 45173.06        | California | 14681.40 |

### 3. Data Understanding

In [3]: `data_startups.shape`

Out[3]: (50, 5)

In [4]: `data_startups.isna().sum()`

Out[4]: R&D Spend            0  
Administration        0  
Marketing Spend        0  
State                    0  
Profit                    0  
dtype: int64

In [5]: `data_startups.dtypes`

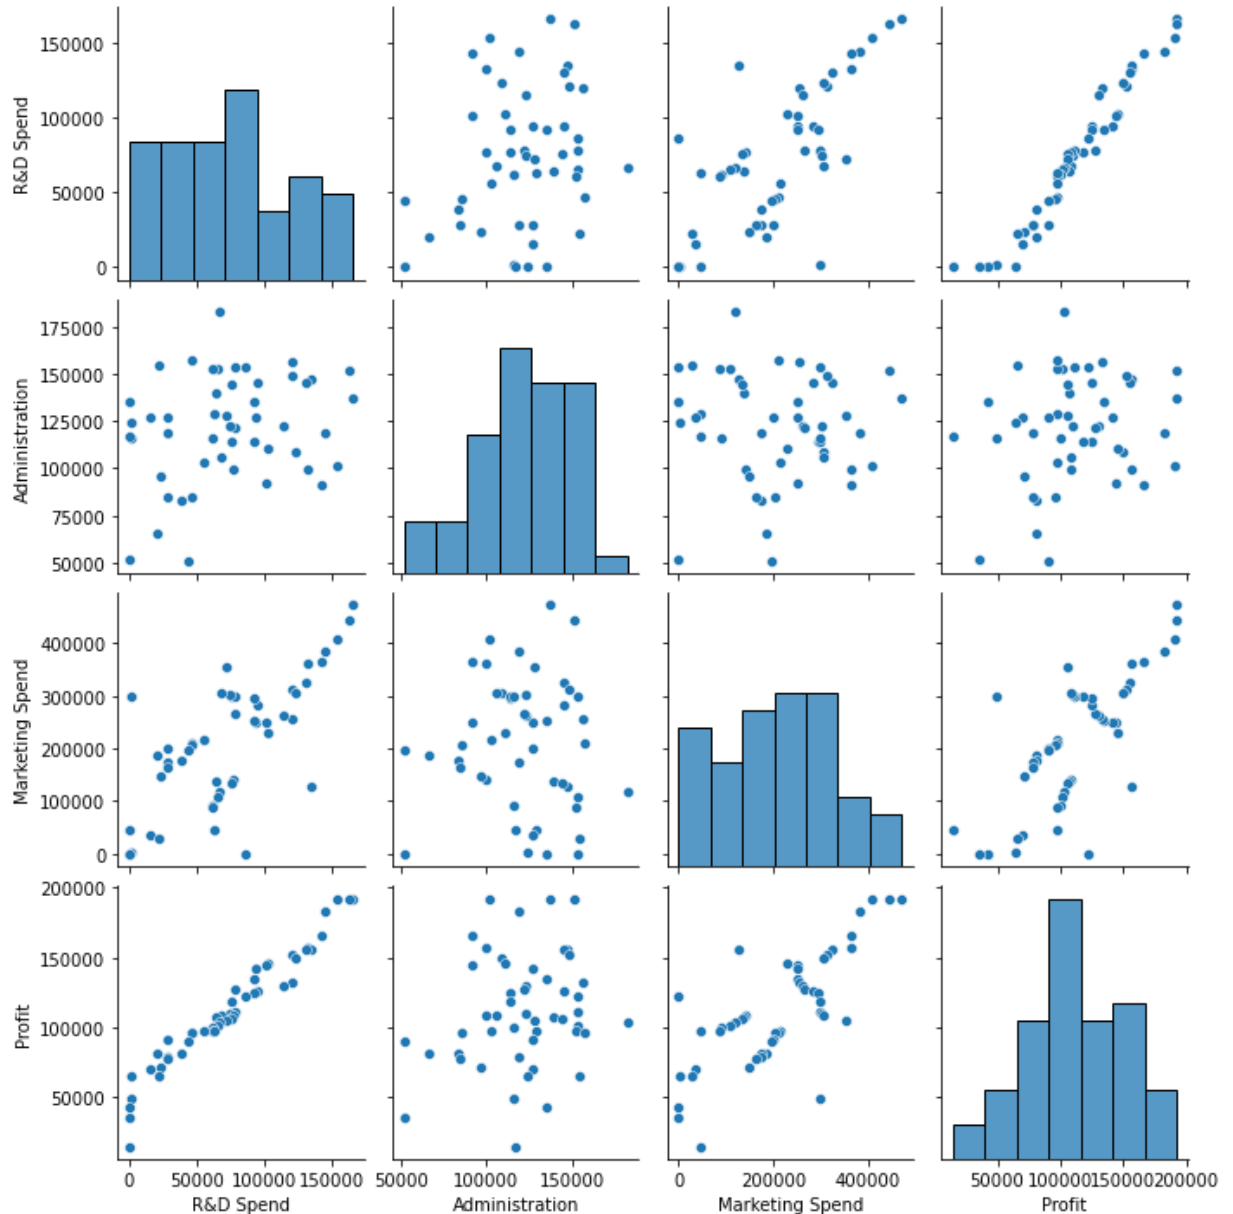
Out[5]: R&D Spend            float64  
Administration        float64  
Marketing Spend        float64  
State                    object  
Profit                    float64  
dtype: object

### 3. Check whether the assumptions is matching or not

### 3.1. Linearity Check

- By using Scatter plot/Pairplot

```
In [6]: sns.pairplot(data_startups)
plt.show()
```



**Observation - Linearity check is failed.**

### 3.2. No Multicollinearity

By using,

Correlation Matrix (or), Variance Inflation Factor(VIF) we can check it.

```
In [7]: corr_martrix = data_startups.corr().round()
corr_martrix
```

```
Out[7]:
```

|                 | R&D Spend | Administration | Marketing Spend | Profit |
|-----------------|-----------|----------------|-----------------|--------|
| R&D Spend       | 1.0       | 0.0            | 1.0             | 1.0    |
| Administration  | 0.0       | 1.0            | -0.0            | 0.0    |
| Marketing Spend | 1.0       | -0.0           | 1.0             | 1.0    |
| Profit          | 1.0       | 0.0            | 1.0             | 1.0    |

```
In [8]: sns.heatmap(data= corr_martrix,annot=True)
plt.show()
```



### 3.3. No AutoRegression

It is satisfied.

### 3.4. Homoscedascity check

This can be done after model training.

### 3.5. Zero Residual Mean

This also can be checked after model training.

## 4.Data Preparation

```
In [3]: del data_startups['State']
```

```
In [4]: data_startups
```

```
Out[4]:
```

|           | <b>R&amp;D Spend</b> | <b>Administration</b> | <b>Marketing Spend</b> | <b>Profit</b> |
|-----------|----------------------|-----------------------|------------------------|---------------|
| <b>0</b>  | 165349.20            | 136897.80             | 471784.10              | 192261.83     |
| <b>1</b>  | 162597.70            | 151377.59             | 443898.53              | 191792.06     |
| <b>2</b>  | 153441.51            | 101145.55             | 407934.54              | 191050.39     |
| <b>3</b>  | 144372.41            | 118671.85             | 383199.62              | 182901.99     |
| <b>4</b>  | 142107.34            | 91391.77              | 366168.42              | 166187.94     |
| <b>5</b>  | 131876.90            | 99814.71              | 362861.36              | 156991.12     |
| <b>6</b>  | 134615.46            | 147198.87             | 127716.82              | 156122.51     |
| <b>7</b>  | 130298.13            | 145530.06             | 323876.68              | 155752.60     |
| <b>8</b>  | 120542.52            | 148718.95             | 311613.29              | 152211.77     |
| <b>9</b>  | 123334.88            | 108679.17             | 304981.62              | 149759.96     |
| <b>10</b> | 101913.08            | 110594.11             | 229160.95              | 146121.95     |
| <b>11</b> | 100671.96            | 91790.61              | 249744.55              | 144259.40     |

```
In [6]: data_startups=data_startups.rename({'R&D Spend':'RDS','Administration':'ADMS','Marketing Spend':'MKTS'})
data_startups
```

```
Out[6]:
```

|    | RDS       | ADMS      | MKTS      | Profit    |
|----|-----------|-----------|-----------|-----------|
| 0  | 165349.20 | 136897.80 | 471784.10 | 192261.83 |
| 1  | 162597.70 | 151377.59 | 443898.53 | 191792.06 |
| 2  | 153441.51 | 101145.55 | 407934.54 | 191050.39 |
| 3  | 144372.41 | 118671.85 | 383199.62 | 182901.99 |
| 4  | 142107.34 | 91391.77  | 366168.42 | 166187.94 |
| 5  | 131876.90 | 99814.71  | 362861.36 | 156991.12 |
| 6  | 134615.46 | 147198.87 | 127716.82 | 156122.51 |
| 7  | 130298.13 | 145530.06 | 323876.68 | 155752.60 |
| 8  | 120542.52 | 148718.95 | 311613.29 | 152211.77 |
| 9  | 123334.88 | 108679.17 | 304981.62 | 149759.96 |
| 10 | 101913.08 | 110594.11 | 229160.95 | 146121.95 |
| 11 | 100671.96 | 91790.61  | 249744.55 | 144259.40 |
| 12 | 93863.75  | 127320.38 | 249839.44 | 141585.52 |
| 13 | 91992.39  | 135495.07 | 252664.93 | 134307.35 |
| 14 | 119943.24 | 156547.42 | 256512.92 | 132602.65 |
| 15 | 114523.61 | 122616.84 | 261776.23 | 129917.04 |
| 16 | 78013.11  | 121597.55 | 264346.06 | 126992.93 |
| 17 | 94657.16  | 145077.58 | 282574.31 | 125370.37 |
| 18 | 91749.16  | 114175.79 | 294919.57 | 124266.90 |
| 19 | 86419.70  | 153514.11 | 0.00      | 122776.86 |
| 20 | 76253.86  | 113867.30 | 298664.47 | 118474.03 |
| 21 | 78389.47  | 153773.43 | 299737.29 | 111313.02 |
| 22 | 73994.56  | 122782.75 | 303319.26 | 110352.25 |
| 23 | 67532.53  | 105751.03 | 304768.73 | 108733.99 |
| 24 | 77044.01  | 99281.34  | 140574.81 | 108552.04 |
| 25 | 64664.71  | 139553.16 | 137962.62 | 107404.34 |
| 26 | 75328.87  | 144135.98 | 134050.07 | 105733.54 |
| 27 | 72107.60  | 127864.55 | 353183.81 | 105008.31 |
| 28 | 66051.52  | 182645.56 | 118148.20 | 103282.38 |
| 29 | 65605.48  | 153032.06 | 107138.38 | 101004.64 |
| 30 | 61994.48  | 115641.28 | 91131.24  | 99937.59  |
| 31 | 61136.38  | 152701.92 | 88218.23  | 97483.56  |
| 32 | 63408.86  | 129219.61 | 46085.25  | 97427.84  |
| 33 | 55493.95  | 103057.49 | 214634.81 | 96778.92  |

|    | RDS      | ADMS      | MKTS      | Profit   |
|----|----------|-----------|-----------|----------|
| 34 | 46426.07 | 157693.92 | 210797.67 | 96712.80 |
| 35 | 46014.02 | 85047.44  | 205517.64 | 96479.51 |
| 36 | 28663.76 | 127056.21 | 201126.82 | 90708.19 |
| 37 | 44069.95 | 51283.14  | 197029.42 | 89949.14 |
| 38 | 20229.59 | 65947.93  | 185265.10 | 81229.06 |
| 39 | 38558.51 | 82982.09  | 174999.30 | 81005.76 |
| 40 | 28754.33 | 118546.05 | 172795.67 | 78239.91 |
| 41 | 27892.92 | 84710.77  | 164470.71 | 77798.83 |
| 42 | 23640.93 | 96189.63  | 148001.11 | 71498.49 |
| 43 | 15505.73 | 127382.30 | 35534.17  | 69758.98 |
| 44 | 22177.74 | 154806.14 | 28334.72  | 65200.33 |
| 45 | 1000.23  | 124153.04 | 1903.93   | 64926.08 |
| 46 | 1315.46  | 115816.21 | 297114.46 | 49490.75 |
| 47 | 0.00     | 135426.92 | 0.00      | 42559.73 |
| 48 | 542.05   | 51743.15  | 0.00      | 35673.41 |
| 49 | 0.00     | 116983.80 | 45173.06  | 14681.40 |

## 5. Model Building | Training | Evaluating using Statsmodels

```
In [29]: import statsmodels.formula.api as smf
model = smf.ols('Profit~RDS+ADMS+MKTS',data= data_startups).fit()
```

```
In [30]: model.params
```

```
Out[30]: Intercept    50122.192990
RDS                0.805715
ADMS               -0.026816
MKTS               0.027228
dtype: float64
```

```
In [31]: model.pvalues
```

```
Out[31]: Intercept    1.057379e-09
RDS                2.634968e-22
ADMS               6.017551e-01
MKTS               1.047168e-01
dtype: float64
```

```
In [32]: model.rsquared,model.rsquared_adj
```

```
Out[32]: (0.9507459940683246, 0.9475337762901719)
```



## Understanding R2 and AdjustedR2

```
In [11]: model_1 = smf.ols('Profit~RDS',data = data_startups).fit()
print('R2 score           : ',round(model_1.rsquared,4))
print('Adjusted R2 score : ',round(model_1.rsquared_adj,4))
print('AIC value         : ',round(model_1.aic,4)) #is an estimator of prediction
print('BIC value         : ',round(model_1.bic,4)) ##is an estimator of prediction
```

```
R2 score           : 0.9465
Adjusted R2 score  : 0.9454
AIC value          : 1058.873
BIC value          : 1062.6971
```

```
In [12]: model_2 = smf.ols('Profit~RDS+ADMS',data = data_startups).fit()
print('R2 score           : ',round(model_2.rsquared,4))
print('Adjusted R2 score : ',round(model_2.rsquared_adj,4))
print('AIC value         : ',round(model_2.aic,4))
print('BIC value         : ',round(model_2.bic,4))
```

```
R2 score           : 0.9478
Adjusted R2 score  : 0.9456
AIC value          : 1059.6637
BIC value          : 1065.3998
```

```
In [13]: model_3 = smf.ols('Profit~RDS+ADMS+MKTS',data = data_startups).fit()
print('R2 score           : ',round(model_3.rsquared,4))
print('Adjusted R2 score : ',round(model_3.rsquared_adj,4))
print('AIC value         : ',round(model_3.aic,4))
print('BIC value         : ',round(model_3.bic,4))
```

```
R2 score           : 0.9507
Adjusted R2 score  : 0.9475
AIC value          : 1058.7715
BIC value          : 1066.4196
```

### Observation:

Always R2 score increases if we increase the number of inputs. And if the input contributes more for the prediction, there will a higher increase in the R2 score.

### To look the complete Summary of the Model

In [14]: `model_1.summary()`

Out[14]: OLS Regression Results

|                          |                  |                            |          |
|--------------------------|------------------|----------------------------|----------|
| <b>Dep. Variable:</b>    | Profit           | <b>R-squared:</b>          | 0.947    |
| <b>Model:</b>            | OLS              | <b>Adj. R-squared:</b>     | 0.945    |
| <b>Method:</b>           | Least Squares    | <b>F-statistic:</b>        | 849.8    |
| <b>Date:</b>             | Fri, 29 Oct 2021 | <b>Prob (F-statistic):</b> | 3.50e-32 |
| <b>Time:</b>             | 22:31:27         | <b>Log-Likelihood:</b>     | -527.44  |
| <b>No. Observations:</b> | 50               | <b>AIC:</b>                | 1059.    |
| <b>Df Residuals:</b>     | 48               | <b>BIC:</b>                | 1063.    |
| <b>Df Model:</b>         | 1                |                            |          |
| <b>Covariance Type:</b>  | nonrobust        |                            |          |

|                  | coef      | std err  | t      | P> t  | [0.025   | 0.975]   |
|------------------|-----------|----------|--------|-------|----------|----------|
| <b>Intercept</b> | 4.903e+04 | 2537.897 | 19.320 | 0.000 | 4.39e+04 | 5.41e+04 |
| <b>RDS</b>       | 0.8543    | 0.029    | 29.151 | 0.000 | 0.795    | 0.913    |

|                       |        |                          |          |
|-----------------------|--------|--------------------------|----------|
| <b>Omnibus:</b>       | 13.727 | <b>Durbin-Watson:</b>    | 1.116    |
| <b>Prob(Omnibus):</b> | 0.001  | <b>Jarque-Bera (JB):</b> | 18.536   |
| <b>Skew:</b>          | -0.911 | <b>Prob(JB):</b>         | 9.44e-05 |
| <b>Kurtosis:</b>      | 5.361  | <b>Cond. No.</b>         | 1.65e+05 |

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 1.65e+05. This might indicate that there are strong multicollinearity or other numerical problems.

In [15]: `model_2.summary()`

Out[15]: OLS Regression Results

|                          |                  |                            |          |
|--------------------------|------------------|----------------------------|----------|
| <b>Dep. Variable:</b>    | Profit           | <b>R-squared:</b>          | 0.948    |
| <b>Model:</b>            | OLS              | <b>Adj. R-squared:</b>     | 0.946    |
| <b>Method:</b>           | Least Squares    | <b>F-statistic:</b>        | 426.8    |
| <b>Date:</b>             | Fri, 29 Oct 2021 | <b>Prob (F-statistic):</b> | 7.29e-31 |
| <b>Time:</b>             | 22:31:29         | <b>Log-Likelihood:</b>     | -526.83  |
| <b>No. Observations:</b> | 50               | <b>AIC:</b>                | 1060.    |
| <b>Df Residuals:</b>     | 47               | <b>BIC:</b>                | 1065.    |
| <b>Df Model:</b>         | 2                |                            |          |
| <b>Covariance Type:</b>  | nonrobust        |                            |          |

|                  | coef      | std err  | t      | P> t  | [0.025   | 0.975]  |
|------------------|-----------|----------|--------|-------|----------|---------|
| <b>Intercept</b> | 5.489e+04 | 6016.718 | 9.122  | 0.000 | 4.28e+04 | 6.7e+04 |
| <b>RDS</b>       | 0.8621    | 0.030    | 28.589 | 0.000 | 0.801    | 0.923   |
| <b>ADMS</b>      | -0.0530   | 0.049    | -1.073 | 0.289 | -0.152   | 0.046   |

|                       |        |                          |          |
|-----------------------|--------|--------------------------|----------|
| <b>Omnibus:</b>       | 14.678 | <b>Durbin-Watson:</b>    | 1.189    |
| <b>Prob(Omnibus):</b> | 0.001  | <b>Jarque-Bera (JB):</b> | 20.449   |
| <b>Skew:</b>          | -0.961 | <b>Prob(JB):</b>         | 3.63e-05 |
| <b>Kurtosis:</b>      | 5.474  | <b>Cond. No.</b>         | 6.65e+05 |

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 6.65e+05. This might indicate that there are strong multicollinearity or other numerical problems.

In [16]: `model_3.summary()`

Out[16]: OLS Regression Results

|                          |                  |                            |          |
|--------------------------|------------------|----------------------------|----------|
| <b>Dep. Variable:</b>    | Profit           | <b>R-squared:</b>          | 0.951    |
| <b>Model:</b>            | OLS              | <b>Adj. R-squared:</b>     | 0.948    |
| <b>Method:</b>           | Least Squares    | <b>F-statistic:</b>        | 296.0    |
| <b>Date:</b>             | Fri, 29 Oct 2021 | <b>Prob (F-statistic):</b> | 4.53e-30 |
| <b>Time:</b>             | 22:31:31         | <b>Log-Likelihood:</b>     | -525.39  |
| <b>No. Observations:</b> | 50               | <b>AIC:</b>                | 1059.    |
| <b>Df Residuals:</b>     | 46               | <b>BIC:</b>                | 1066.    |
| <b>Df Model:</b>         | 3                |                            |          |
| <b>Covariance Type:</b>  | nonrobust        |                            |          |

|                  | coef      | std err  | t      | P> t  | [0.025   | 0.975]   |
|------------------|-----------|----------|--------|-------|----------|----------|
| <b>Intercept</b> | 5.012e+04 | 6572.353 | 7.626  | 0.000 | 3.69e+04 | 6.34e+04 |
| <b>RDS</b>       | 0.8057    | 0.045    | 17.846 | 0.000 | 0.715    | 0.897    |
| <b>ADMS</b>      | -0.0268   | 0.051    | -0.526 | 0.602 | -0.130   | 0.076    |
| <b>MKTS</b>      | 0.0272    | 0.016    | 1.655  | 0.105 | -0.006   | 0.060    |

|                       |        |                          |          |
|-----------------------|--------|--------------------------|----------|
| <b>Omnibus:</b>       | 14.838 | <b>Durbin-Watson:</b>    | 1.282    |
| <b>Prob(Omnibus):</b> | 0.001  | <b>Jarque-Bera (JB):</b> | 21.442   |
| <b>Skew:</b>          | -0.949 | <b>Prob(JB):</b>         | 2.21e-05 |
| <b>Kurtosis:</b>      | 5.586  | <b>Cond. No.</b>         | 1.40e+06 |

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.4e+06. This might indicate that there are strong multicollinearity or other numerical problems.

## 6. Model Validation

- Two Techniques:
  1. Collinearity Check &
  2. Residual Analysis

```
In [17]: rsq_r=smf.ols("RDS~ADMS+MKTS",data=data_startups).fit().rsquared
vif_r=1/(1-rsq_r)

rsq_a=smf.ols("ADMS~RDS+MKTS",data=data_startups).fit().rsquared
vif_a=1/(1-rsq_a)

rsq_m=smf.ols("MKTS~RDS+ADMS",data=data_startups).fit().rsquared
vif_m=1/(1-rsq_m)

# Putting the values in Dataframe format
d1={'Variables':['RDS','ADMS','MKTS'],'Vif':[vif_r,vif_a,vif_m]}
Vif_df=pd.DataFrame(d1)
Vif_df
```

```
Out[17]:
```

|   | Variables | Vif      |
|---|-----------|----------|
| 0 | RDS       | 2.468903 |
| 1 | ADMS      | 1.175091 |
| 2 | MKTS      | 2.326773 |


observations-None variable has VIF>20, No Collinearity, so consider all variables in Regression equation

## Test for Normality of Residuals (Q-Q Plot)

```
In [33]: import statsmodels.api as sm

# 2) Residual Analysis
# Test for Normality of Residuals (Q-Q Plot) using residual model (model.resid)

sm.qqplot(model.resid,line='q')
plt.title("Normal Q-Q plot of residuals")
plt.show()
```



```
In [34]: list(np.where(model.resid<-30000))
```


```
Out[34]: [array([49], dtype=int64)]
```

## Residual Plot for Homoscedasticity

```
In [23]: # Test for Homoscedasticity or Heteroscedasticity (plotting model's standardized
```

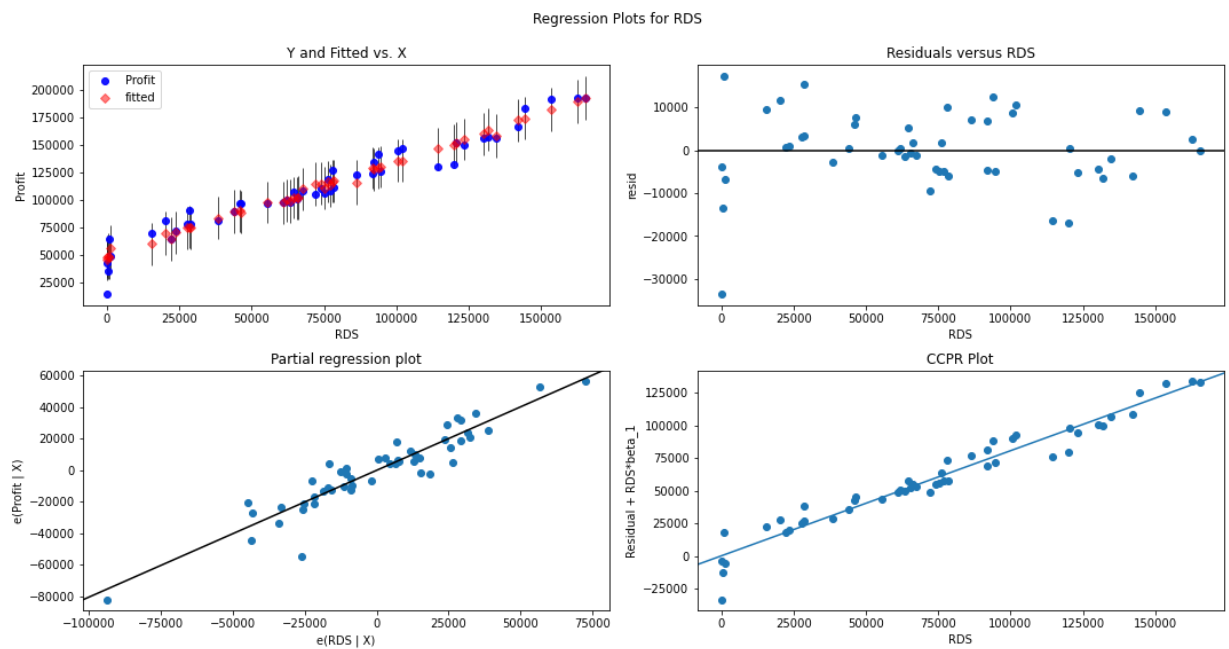
```
def standard_values(vals) :  
    return (vals-vals.mean())/vals.std() # User defined  $z = (x - \mu)/\sigma$ 
```

```
In [35]: plt.scatter(standard_values(model.fittedvalues),standard_values(model_1.resid))  
plt.title('Residual Plot')  
plt.xlabel('standardized fitted values')  
plt.ylabel('standardized residual values')  
plt.show()
```

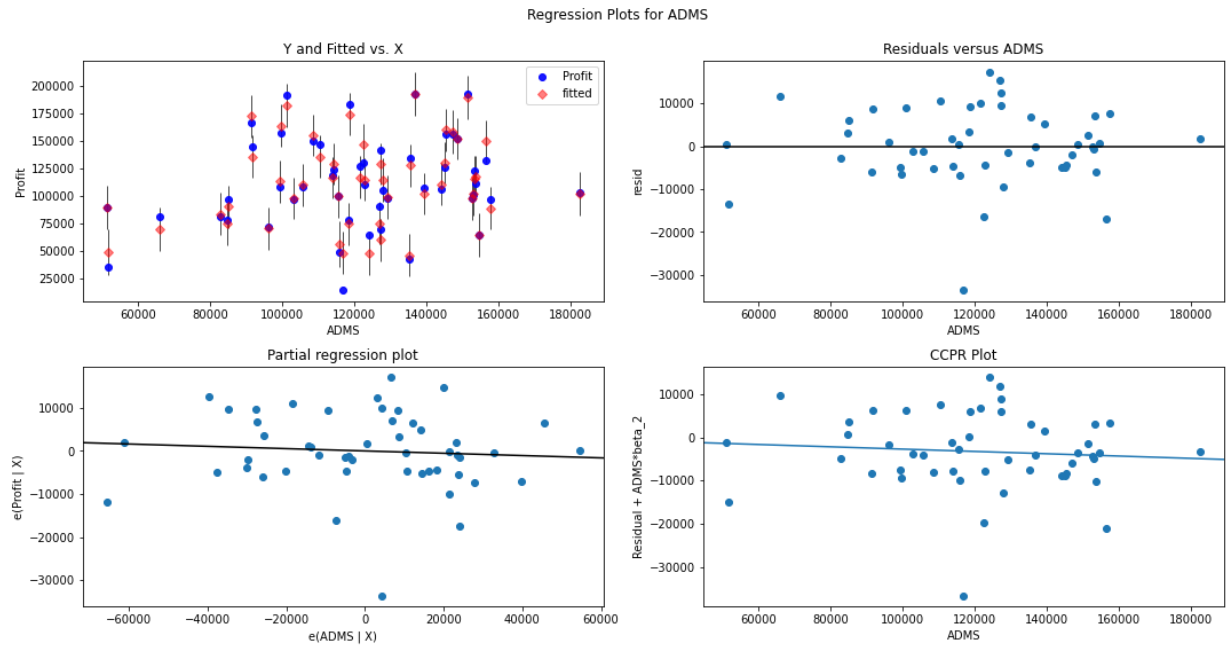


## Residual Vs Regressors

```
In [36]: fig=plt.figure(figsize=(15,8))
sm.graphics.plot_regress_exog(model,'RDS',fig=fig)
plt.show()
```

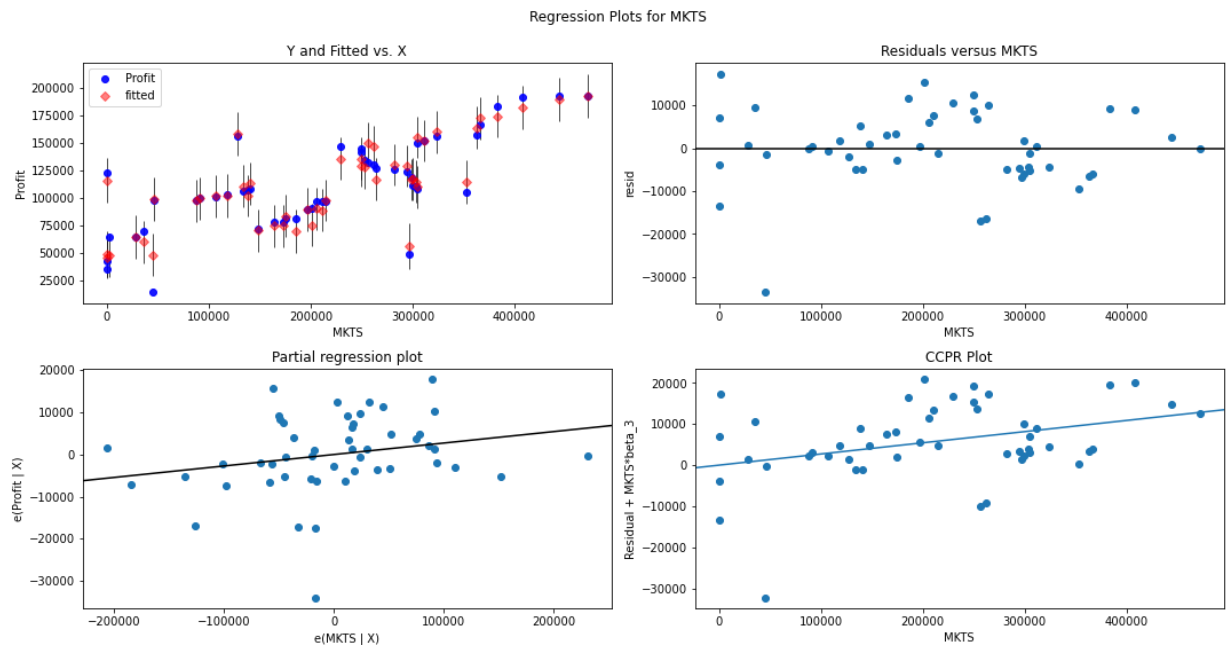


```
In [38]: fig=plt.figure(figsize=(15,8))
sm.graphics.plot_regress_exog(model,'ADMS',fig=fig)
plt.show()
```





```
In [39]: fig=plt.figure(figsize=(15,8))
sm.graphics.plot_regress_exog(model,'MKTS',fig=fig)
plt.show()
```



## Model Deletion Diagnostics


## Detecting Influencers/Outliers

### Cook's Distance

```
In [40]: # 1. Cook's Distance: If Cook's distance > 1, then it's an outlier
# Get influencers using cook's distance
(c,_)=model.get_influence().cooks_distance
c
```

```
Out[40]: array([3.21825244e-05, 3.27591036e-03, 3.23842699e-02, 2.17206555e-02,
1.44833032e-02, 1.17158463e-02, 2.91766303e-03, 3.56513444e-03,
4.04303948e-05, 4.86758017e-03, 1.51064757e-02, 1.63564959e-02,
1.15516625e-02, 4.01422811e-03, 6.12934253e-02, 3.40013448e-02,
8.33556413e-03, 3.30534399e-03, 2.16819303e-03, 4.07440577e-02,
4.25137222e-04, 1.09844352e-02, 2.91768000e-03, 2.76030254e-04,
5.04643588e-03, 3.00074623e-03, 3.41957068e-03, 2.98396413e-02,
1.31590664e-03, 1.25992620e-04, 4.18505125e-05, 9.27434786e-06,
7.08656521e-04, 1.28122674e-04, 2.09815032e-02, 6.69508674e-03,
5.55314705e-02, 6.55050578e-05, 5.61547311e-02, 1.54279607e-03,
1.84850929e-03, 1.97578066e-03, 1.36089280e-04, 2.05553171e-02,
1.23156041e-04, 9.03234206e-02, 5.45303387e-02, 5.33885616e-03,
1.90527441e-01, 2.88082293e-01])
```

```
In [42]: # Plot the influencers using the stem plot
fig=plt.figure(figsize=(20,7))
plt.stem(np.arange(len(data_startups)),np.round(c,5))
plt.xlabel('Row Index')
plt.ylabel('Cooks Distance')
plt.show()
```




```
In [44]: # Index and value of influencer where C>0.5
np.argmax(c) , np.max(c)
```

```
Out[44]: (49, 0.28808229275432584)
```

**# 2. Leverage Value using High Influence Points : Points beyond Leverage\_cutoff value are influencers**

```
In [46]: from statsmodels.graphics.regressionplots import influence_plot

influence_plot(model)
plt.show()
```



```
In [47]: # Leverage Cutoff Value = 3*(k+1)/n ; k = no.of features/columns & n = no. of data points
k=data_startups.shape[1]
n=data_startups.shape[0]
leverage_cutoff = (3*(k+1))/n
leverage_cutoff
```

Out[47]: 0.3

```
In [48]: data_startups[data_startups.index.isin([49])]
```

Out[48]:

|    | RDS | ADMS     | MKTS     | Profit  |
|----|-----|----------|----------|---------|
| 49 | 0.0 | 116983.8 | 45173.06 | 14681.4 |

```
In [49]: data_startups.head()
```

Out[49]:

|   | RDS       | ADMS      | MKTS      | Profit    |
|---|-----------|-----------|-----------|-----------|
| 0 | 165349.20 | 136897.80 | 471784.10 | 192261.83 |
| 1 | 162597.70 | 151377.59 | 443898.53 | 191792.06 |
| 2 | 153441.51 | 101145.55 | 407934.54 | 191050.39 |
| 3 | 144372.41 | 118671.85 | 383199.62 | 182901.99 |
| 4 | 142107.34 | 91391.77  | 366168.42 | 166187.94 |

## Improving the model

**Discard the data points which are influencers and reassign the row number (`reset_index(drop=True)`)**

```
In [51]: data_2=data_startups.drop(data_startups.index[[49]],axis=0).reset_index(drop=True)
data_2
```

```
Out[51]:
```

|    | RDS       | ADMS      | MKTS      | Profit    |
|----|-----------|-----------|-----------|-----------|
| 0  | 165349.20 | 136897.80 | 471784.10 | 192261.83 |
| 1  | 162597.70 | 151377.59 | 443898.53 | 191792.06 |
| 2  | 153441.51 | 101145.55 | 407934.54 | 191050.39 |
| 3  | 144372.41 | 118671.85 | 383199.62 | 182901.99 |
| 4  | 142107.34 | 91391.77  | 366168.42 | 166187.94 |
| 5  | 131876.90 | 99814.71  | 362861.36 | 156991.12 |
| 6  | 134615.46 | 147198.87 | 127716.82 | 156122.51 |
| 7  | 130298.13 | 145530.06 | 323876.68 | 155752.60 |
| 8  | 120542.52 | 148718.95 | 311613.29 | 152211.77 |
| 9  | 123334.88 | 108679.17 | 304981.62 | 149759.96 |
| 10 | 101913.08 | 110594.11 | 229160.95 | 146121.95 |
| 11 | 100671.96 | 91790.61  | 249744.55 | 144259.40 |
| 12 | 93863.75  | 127320.38 | 249839.44 | 141585.52 |
| 13 | 91992.39  | 135495.07 | 252664.93 | 134307.35 |
| 14 | 119943.24 | 156547.42 | 256512.92 | 132602.65 |
| 15 | 114523.61 | 122616.84 | 261776.23 | 129917.04 |
| 16 | 78013.11  | 121597.55 | 264346.06 | 126992.93 |
| 17 | 94657.16  | 145077.58 | 282574.31 | 125370.37 |
| 18 | 91749.16  | 114175.79 | 294919.57 | 124266.90 |
| 19 | 86419.70  | 153514.11 | 0.00      | 122776.86 |
| 20 | 76253.86  | 113867.30 | 298664.47 | 118474.03 |
| 21 | 78389.47  | 153773.43 | 299737.29 | 111313.02 |
| 22 | 73994.56  | 122782.75 | 303319.26 | 110352.25 |
| 23 | 67532.53  | 105751.03 | 304768.73 | 108733.99 |
| 24 | 77044.01  | 99281.34  | 140574.81 | 108552.04 |
| 25 | 64664.71  | 139553.16 | 137962.62 | 107404.34 |
| 26 | 75328.87  | 144135.98 | 134050.07 | 105733.54 |
| 27 | 72107.60  | 127864.55 | 353183.81 | 105008.31 |
| 28 | 66051.52  | 182645.56 | 118148.20 | 103282.38 |
| 29 | 65605.48  | 153032.06 | 107138.38 | 101004.64 |
| 30 | 61994.48  | 115641.28 | 91131.24  | 99937.59  |
| 31 | 61136.38  | 152701.92 | 88218.23  | 97483.56  |
| 32 | 63408.86  | 129219.61 | 46085.25  | 97427.84  |
| 33 | 55493.95  | 103057.49 | 214634.81 | 96778.92  |

|    | RDS      | ADMS      | MKTS      | Profit   |
|----|----------|-----------|-----------|----------|
| 34 | 46426.07 | 157693.92 | 210797.67 | 96712.80 |
| 35 | 46014.02 | 85047.44  | 205517.64 | 96479.51 |
| 36 | 28663.76 | 127056.21 | 201126.82 | 90708.19 |
| 37 | 44069.95 | 51283.14  | 197029.42 | 89949.14 |
| 38 | 20229.59 | 65947.93  | 185265.10 | 81229.06 |
| 39 | 38558.51 | 82982.09  | 174999.30 | 81005.76 |
| 40 | 28754.33 | 118546.05 | 172795.67 | 78239.91 |
| 41 | 27892.92 | 84710.77  | 164470.71 | 77798.83 |
| 42 | 23640.93 | 96189.63  | 148001.11 | 71498.49 |
| 43 | 15505.73 | 127382.30 | 35534.17  | 69758.98 |
| 44 | 22177.74 | 154806.14 | 28334.72  | 65200.33 |
| 45 | 1000.23  | 124153.04 | 1903.93   | 64926.08 |
| 46 | 1315.46  | 115816.21 | 297114.46 | 49490.75 |
| 47 | 0.00     | 135426.92 | 0.00      | 42559.73 |
| 48 | 542.05   | 51743.15  | 0.00      | 35673.41 |

## Model Deletion Diagnostics and Final Model

```
In [52]: model_new=smf.ols("Profit~RDS+ADMS+MKTS",data=data_2).fit()
```

```
In [53]: while model_new.rsquared < 0.99:
          for c in [np.max(c)>1]:
              model_new=smf.ols("Profit~RDS+ADMS+MKTS",data=data_2).fit()
              (c,_)=model_new.get_influence().cooks_distance
              c
              np.argmax(c) , np.max(c)
              data_2=data_2.drop(data_2.index[[np.argmax(c)]],axis=0).reset_index(drop=
              data_2
          else:
              final_model=smf.ols("Profit~RDS+ADMS+MKTS",data=data_2).fit()
              final_model.rsquared , final_model.aic
              print("Thus model accuracy is improved to",final_model.rsquared)
```

```
Thus model accuracy is improved to 0.9626766170294073
Thus model accuracy is improved to 0.9614129113440602
Thus model accuracy is improved to 0.962593650298269
Thus model accuracy is improved to 0.9638487279209413
Thus model accuracy is improved to 0.9663901957918793
Thus model accuracy is improved to 0.9706076169779905
Thus model accuracy is improved to 0.9727840588916423
Thus model accuracy is improved to 0.9734292907181952
Thus model accuracy is improved to 0.9785801571833451
Thus model accuracy is improved to 0.9777383743090916
Thus model accuracy is improved to 0.9790510088977512
Thus model accuracy is improved to 0.9790004461890552
Thus model accuracy is improved to 0.9807878666153609
Thus model accuracy is improved to 0.9838299343609735
Thus model accuracy is improved to 0.983114992639277
Thus model accuracy is improved to 0.9833768520972176
Thus model accuracy is improved to 0.9878892536376698
Thus model accuracy is improved to 0.9877191935547199
Thus model accuracy is improved to 0.9858356627471713
Thus model accuracy is improved to 0.9874766829880098
Thus model accuracy is improved to 0.9906666289527223
Thus model accuracy is improved to 0.9882757054424702
```

```
In [54]: final_model.rsquared
```

```
Out[54]: 0.9882757054424702
```

In [55]: data\_2

Out[55]:

|           | <b>RDS</b> | <b>ADMS</b> | <b>MKTS</b> | <b>Profit</b> |
|-----------|------------|-------------|-------------|---------------|
| <b>0</b>  | 142107.34  | 91391.77    | 366168.42   | 166187.94     |
| <b>1</b>  | 131876.90  | 99814.71    | 362861.36   | 156991.12     |
| <b>2</b>  | 130298.13  | 145530.06   | 323876.68   | 155752.60     |
| <b>3</b>  | 120542.52  | 148718.95   | 311613.29   | 152211.77     |
| <b>4</b>  | 123334.88  | 108679.17   | 304981.62   | 149759.96     |
| <b>5</b>  | 91992.39   | 135495.07   | 252664.93   | 134307.35     |
| <b>6</b>  | 94657.16   | 145077.58   | 282574.31   | 125370.37     |
| <b>7</b>  | 91749.16   | 114175.79   | 294919.57   | 124266.90     |
| <b>8</b>  | 76253.86   | 113867.30   | 298664.47   | 118474.03     |
| <b>9</b>  | 67532.53   | 105751.03   | 304768.73   | 108733.99     |
| <b>10</b> | 77044.01   | 99281.34    | 140574.81   | 108552.04     |
| <b>11</b> | 64664.71   | 139553.16   | 137962.62   | 107404.34     |
| <b>12</b> | 75328.87   | 144135.98   | 134050.07   | 105733.54     |
| <b>13</b> | 66051.52   | 182645.56   | 118148.20   | 103282.38     |
| <b>14</b> | 65605.48   | 153032.06   | 107138.38   | 101004.64     |
| <b>15</b> | 61994.48   | 115641.28   | 91131.24    | 99937.59      |
| <b>16</b> | 61136.38   | 152701.92   | 88218.23    | 97483.56      |
| <b>17</b> | 63408.86   | 129219.61   | 46085.25    | 97427.84      |
| <b>18</b> | 55493.95   | 103057.49   | 214634.81   | 96778.92      |
| <b>19</b> | 46426.07   | 157693.92   | 210797.67   | 96712.80      |
| <b>20</b> | 46014.02   | 85047.44    | 205517.64   | 96479.51      |
| <b>21</b> | 44069.95   | 51283.14    | 197029.42   | 89949.14      |
| <b>22</b> | 38558.51   | 82982.09    | 174999.30   | 81005.76      |
| <b>23</b> | 28754.33   | 118546.05   | 172795.67   | 78239.91      |
| <b>24</b> | 27892.92   | 84710.77    | 164470.71   | 77798.83      |
| <b>25</b> | 23640.93   | 96189.63    | 148001.11   | 71498.49      |
| <b>26</b> | 22177.74   | 154806.14   | 28334.72    | 65200.33      |

## Model Predictions



```
In [56]: # say New data for prediction is
new_data=pd.DataFrame({'RDS':70000,"ADMS":90000,"MKTS":140000},index=[0])
new_data
```

```
Out[56]:
```

|   | RDS   | ADMS  | MKTS   |
|---|-------|-------|--------|
| 0 | 70000 | 90000 | 140000 |

```
In [57]: final_model.predict(new_data)
```

```
Out[57]: 0      104858.729408
dtype: float64
```

```
In [58]: # Automatic Prediction of Price with 90.02% accuracy
pred_y=final_model.predict(data_2)
pred_y
```

```
Out[58]: 0      165589.539700
1      158552.826483
2      156789.000710
3      149524.698853
4      150122.356712
5      126598.769555
6      130104.785747
7      127878.387928
8      117298.757074
9      111329.242429
10     110009.916133
11     102331.717613
12     109661.804131
13     103462.767086
14     101874.612012
15      97655.794577
16      97872.919535
17      96858.382686
18      98654.449007
19      93583.600868
20      91186.568204
21      88571.938968
22      84521.312916
23      78528.002935
24      76670.262623
25      73237.524757
26      68075.710756
dtype: float64
```

**table containing R<sup>2</sup> value for each prepared model**

```
In [59]: d2={'Prep_Models':['Model','Final_Model'],'Rsquared':[model.rsquared,final_model.rsquared]}
table=pd.DataFrame(d2)
table
```

```
Out[59]:
```

|   | Prep_Models | Rsquared |
|---|-------------|----------|
| 0 | Model       | 0.950746 |
| 1 | Final_Model | 0.988276 |